


Feedback — VIII. Neural Networks: Representation

[Help](#)

You submitted this quiz on **Mon 11 Nov 2013 5:41 PM PST**. You got a score of **3.75** out of **5.00**. You can [attempt again](#) in 10 minutes.

Question 1


Consider the following neural network which takes two binary-valued inputs $x_1, x_2 \in \{0, 1\}$ and outputs $h_{\Theta}(x)$. Which of the following logical functions does it (approximately) compute?



Your Answer	Score	Explanation
<input checked="" type="radio"/> OR	✓ 1.00	This network will output approximately 1 when either input is 1.
<input type="radio"/> XOR (exclusive OR)		
<input type="radio"/> AND		
<input type="radio"/> NAND (meaning "NOT AND")		
Total	1.00 / 1.00	

Question 2

Consider the neural network given below. Which of the following equations correctly computes the activation $a_1^{(3)}$? Note: $g(z)$ is the sigmoid activation function.



Your Answer

Score

Explanation



$$a_1^{(3)} = g(\Theta_{1,0}^{(2)} a_0^{(1)} + \Theta_{1,1}^{(2)} a_1^{(1)} + \Theta_{1,2}^{(2)} a_2^{(1)})$$



$$a_1^{(3)} = g(\Theta_{2,0}^{(2)} a_0^{(2)} + \Theta_{2,1}^{(2)} a_1^{(2)} + \Theta_{2,2}^{(2)} a_2^{(2)})$$



$$a_1^{(3)} = g(\Theta_{1,0}^{(2)} a_0^{(2)} + \Theta_{1,1}^{(2)} a_1^{(2)} + \Theta_{1,2}^{(2)} a_2^{(2)})$$



1.00

This correctly uses the first row of $\Theta^{(2)}$ and includes the "+1" term of $a_0^{(2)}$.



$$a_1^{(3)} = g(\Theta_{1,0}^{(1)} a_0^{(2)} + \Theta_{1,1}^{(1)} a_1^{(2)} + \Theta_{1,2}^{(1)} a_2^{(2)})$$


Total

1.00 /

1.00

Question 3

You have the following neural network:



You'd like to compute the activations of the hidden layer $a^{(2)} \in \mathbb{R}^3$. One way to do so is the following Octave code:

```
% Theta1 is Theta with superscript "(1)" from lecture
% ie, the matrix of parameters for the mapping from layer 1 (input) to layer 2
% Theta1 has size 3x3
% Assume 'sigmoid' is a built-in function to compute 1 / (1 + exp(-z))

a2 = zeros (3, 1);
for i = 1:3
    for j = 1:3
        a2(i) = a2(i) + x(j) * Theta1(i, j);
    end
    a2(i) = sigmoid (a2(i));
end
```

You want to have a vectorized implementation of this (i.e., one that does not use for loops).

Which of the following implementations correctly compute $a^{(2)}$? Check all that apply.

Your Answer	Score	Explanation
<input type="checkbox"/> <pre>a2 = sigmoid (x * Theta1);</pre>	<input checked="" type="checkbox"/> 0.25	The order of the multiplication is important, this will not work as x is a vector of size 3×1 while Θ_1 is a matrix of size 3×3 .
<input type="checkbox"/> <pre>z = sigmoid(x); a2 = Theta1 * z ;</pre>	<input checked="" type="checkbox"/> 0.25	You should apply the sigmoid function after multiplying with $\Theta^{(1)}$, not before.
<input checked="" type="checkbox"/> <pre>a2 = sigmoid (Theta1 * x);</pre>	<input checked="" type="checkbox"/> 0.25	In the lecture's notation, $a^{(2)} = g(\Theta^{(1)}x)$, so this version computes it directly, as the sigmoid function will act element-wise.



0.25

a2 = sigmoid (
Theta2 * x);

$\Theta^{(2)}$ specifies the parameters from the second to third layers,
not first to second.


Total

1.00 /
1.00

Question 4

You are using the neural network pictured below and have learned the parameters

$\Theta^{(1)} = \begin{bmatrix} 1 & 0.5 & 1.9 \\ 1 & 1.2 & 2.7 \end{bmatrix}$ (used to compute $a^{(2)}$) and $\Theta^{(2)} = \begin{bmatrix} 1 & -0.2 & -1.7 \end{bmatrix}$ (used to compute $a^{(3)}$ as a function of $a^{(2)}$). Suppose you swap the parameters for the first hidden layer between its two units so $\Theta^{(1)} = \begin{bmatrix} 1 & 1.2 & 2.7 \\ 1 & 0.5 & 1.9 \end{bmatrix}$ and also swap the output layer so $\Theta^{(2)} = \begin{bmatrix} 1 & -1.7 & -0.2 \end{bmatrix}$. How will this change the value of the output $h_{\Theta}(x)$?

**Your Answer****Score****Explanation**

☐ It will stay the same.

☒ Insufficient information to tell: it may increase or decrease.



0.00

Swapping $\Theta^{(1)}$ swaps the hidden layers output $a^{(2)}$. But the swap of $\Theta^{(2)}$ cancels out the change, so we can be certain the output will remain unchanged.

☐ It will increase.

☐ It will decrease

Total	0.00 /
	1.00

Question 5

Which of the following statements are true? Check all that apply.

Your Answer	Score	Explanation
<input checked="" type="checkbox"/> Any logical function over binary-valued (0 or 1) inputs x_1 and x_2 can be (approximately) represented using some neural network.	<input checked="" type="checkbox"/> 0.25	Since we can build the basic AND, OR, and NOT functions with a two layer network, we can (approximately) represent any logical function by composing these basic functions over multiple layers.
<input checked="" type="checkbox"/> The activation values of the hidden units in a neural network, with the sigmoid activation function applied at every layer, are always in the range (0, 1).	<input checked="" type="checkbox"/> 0.25	The activation function $g(z) = \frac{1}{1+\exp(-z)}$ has a range of (0, 1).
<input checked="" type="checkbox"/> If a neural network is overfitting the data, one solution would be to decrease the regularization parameter λ .	<input checked="" type="checkbox"/> 0.25	A smaller value of λ allows the model to more closely fit the training data, thereby increasing the chances of overfitting.
<input checked="" type="checkbox"/> A two layer (one input layer, one output layer; no hidden layer) neural network can represent the XOR function.	<input type="checkbox"/> 0.00	We must compose multiple logical operations by using a hidden layer to represent the XOR function.
Total	0.75 /	1.00

